

```
In [ ]: ### Imports
import mrmr
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.preprocessing import PolynomialFeatures
from sklearn.feature_selection import RFECV
from sklearn.linear_model import LogisticRegression
```

```
In [ ]: ### Import data
data = pd.read_csv("train.csv")

# Get general info
print(data.info(), "\n\n\n")
print(data.describe(), "\n\n\n")
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1340 entries, 0 to 1339
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   EmployeeID                           1340 non-null   int64
1   Age                                   1340 non-null   int64
2   Attrition                            1340 non-null   object
3   BusinessTravel                       1340 non-null   object
4   DailyRate                            1340 non-null   int64
5   Department                           1340 non-null   object
6   DistanceFromHome                     1340 non-null   int64
7   Education                             1340 non-null   int64
8   EducationField                       1340 non-null   object
9   EmployeeCount                        1340 non-null   int64
10  EnvironmentSatisfaction               1340 non-null   int64
11  Gender                               1340 non-null   object
12  HourlyRate                           1340 non-null   int64
13  JobInvolvement                       1340 non-null   int64
14  JobLevel                             1340 non-null   int64
15  JobRole                              1340 non-null   object
16  JobSatisfaction                      1340 non-null   int64
17  MaritalStatus                       1340 non-null   object
18  MonthlyIncome                       1340 non-null   int64
19  MonthlyRate                          1340 non-null   int64
20  NumCompaniesWorked                  1340 non-null   int64
21  Over18                              1340 non-null   object
22  OverTime                             1340 non-null   object
23  PercentSalaryHike                   1340 non-null   int64
24  PerformanceRating                   1340 non-null   int64
25  RelationshipSatisfaction             1340 non-null   int64
26  StandardHours                       1340 non-null   int64
27  Shift                               1340 non-null   int64
28  TotalWorkingYears                   1340 non-null   int64
29  TrainingTimesLastYear               1340 non-null   int64
30  WorkLifeBalance                     1340 non-null   int64
31  YearsAtCompany                      1340 non-null   int64
32  YearsInCurrentRole                  1340 non-null   int64
33  YearsSinceLastPromotion              1340 non-null   int64
34  YearsWithCurrManager                 1340 non-null   int64
dtypes: int64(26), object(9)
memory usage: 366.5+ KB
None
```

| | EmployeeID | Age | DailyRate | DistanceFromHome | Educatio |
|-------|--------------|-------------|-------------|------------------|-------------|
| n \ | | | | | |
| count | 1.340000e+03 | 1340.000000 | 1340.000000 | 1340.000000 | 1340.000000 |
| 0 | | | | | |
| mean | 1.460265e+06 | 36.580597 | 799.197761 | 9.193284 | 2.92462 |
| 7 | | | | | |
| std | 2.494821e+05 | 9.013072 | 399.333256 | 8.141621 | 1.03608 |
| 8 | | | | | |
| min | 1.025177e+06 | 18.000000 | 102.000000 | 1.000000 | 1.00000 |
| 0 | | | | | |

| | | | | | |
|----------|--------------|-----------|-------------|-----------|----------|
| 25% 0 | 1.237599e+06 | 30.000000 | 465.000000 | 2.000000 | 2.000000 |
| 50% 0 | 1.469862e+06 | 35.000000 | 796.000000 | 7.000000 | 3.000000 |
| 75% 0 | 1.670131e+06 | 42.000000 | 1153.000000 | 14.000000 | 4.000000 |
| max 0 | 1.886378e+06 | 60.000000 | 1499.000000 | 29.000000 | 5.000000 |

| | EmployeeCount | EnvironmentSatisfaction | HourlyRate | JobInvolvement |
|-------|---------------|-------------------------|-------------|----------------|
| \ | | | | |
| count | 1340.0 | 1340.000000 | 1340.000000 | 1340.000000 |
| mean | 1.0 | 2.709701 | 65.559701 | 2.717910 |
| std | 0.0 | 1.099961 | 20.335025 | 0.717523 |
| min | 1.0 | 1.000000 | 30.000000 | 1.000000 |
| 25% | 1.0 | 2.000000 | 48.000000 | 2.000000 |
| 50% | 1.0 | 3.000000 | 65.000000 | 3.000000 |
| 75% | 1.0 | 4.000000 | 83.000000 | 3.000000 |
| max | 1.0 | 4.000000 | 100.000000 | 4.000000 |

| | JobLevel | ... | RelationshipSatisfaction | StandardHours | Shi |
|-------------|-------------|-----|--------------------------|---------------|-----------|
| ft \ | | | | | |
| count 00 | 1340.000000 | ... | 1340.000000 | 1340.0 | 1340.0000 |
| mean 09 | 2.051493 | ... | 2.700000 | 80.0 | 0.8082 |
| std 51 | 1.104491 | ... | 1.079858 | 0.0 | 0.8562 |
| min 00 | 1.000000 | ... | 1.000000 | 80.0 | 0.0000 |
| 25% 00 | 1.000000 | ... | 2.000000 | 80.0 | 0.0000 |
| 50% 00 | 2.000000 | ... | 3.000000 | 80.0 | 1.0000 |
| 75% 00 | 3.000000 | ... | 4.000000 | 80.0 | 1.0000 |
| max 00 | 5.000000 | ... | 4.000000 | 80.0 | 3.0000 |

| | TotalWorkingYears | TrainingTimesLastYear | WorkLifeBalance | \ |
|-------|-------------------|-----------------------|-----------------|---|
| count | 1340.000000 | 1340.000000 | 1340.000000 | |
| mean | 11.222388 | 2.785821 | 2.771642 | |
| std | 7.696043 | 1.263473 | 0.700007 | |
| min | 0.000000 | 0.000000 | 1.000000 | |
| 25% | 6.000000 | 2.000000 | 2.000000 | |
| 50% | 10.000000 | 3.000000 | 3.000000 | |
| 75% | 15.000000 | 3.000000 | 3.000000 | |
| max | 40.000000 | 6.000000 | 4.000000 | |

| | YearsAtCompany | YearsInCurrentRole | YearsSinceLastPromotion | \ |
|-------|----------------|--------------------|-------------------------|---|
| count | 1340.000000 | 1340.000000 | 1340.000000 | |
| mean | 7.070149 | 4.272388 | 2.175373 | |
| std | 6.039663 | 3.677798 | 3.222376 | |
| min | 0.000000 | 0.000000 | 0.000000 | |
| 25% | 3.000000 | 2.000000 | 0.000000 | |
| 50% | 5.000000 | 3.000000 | 1.000000 | |

| | | | |
|-----|-----------|-----------|-----------|
| 75% | 10.000000 | 7.000000 | 3.000000 |
| max | 40.000000 | 18.000000 | 15.000000 |

```

YearsWithCurrManager
count    1340.000000
mean      4.167164
std       3.581605
min       0.000000
25%       2.000000
50%       3.000000
75%       7.000000
max       17.000000


```


[8 rows x 26 columns]


```

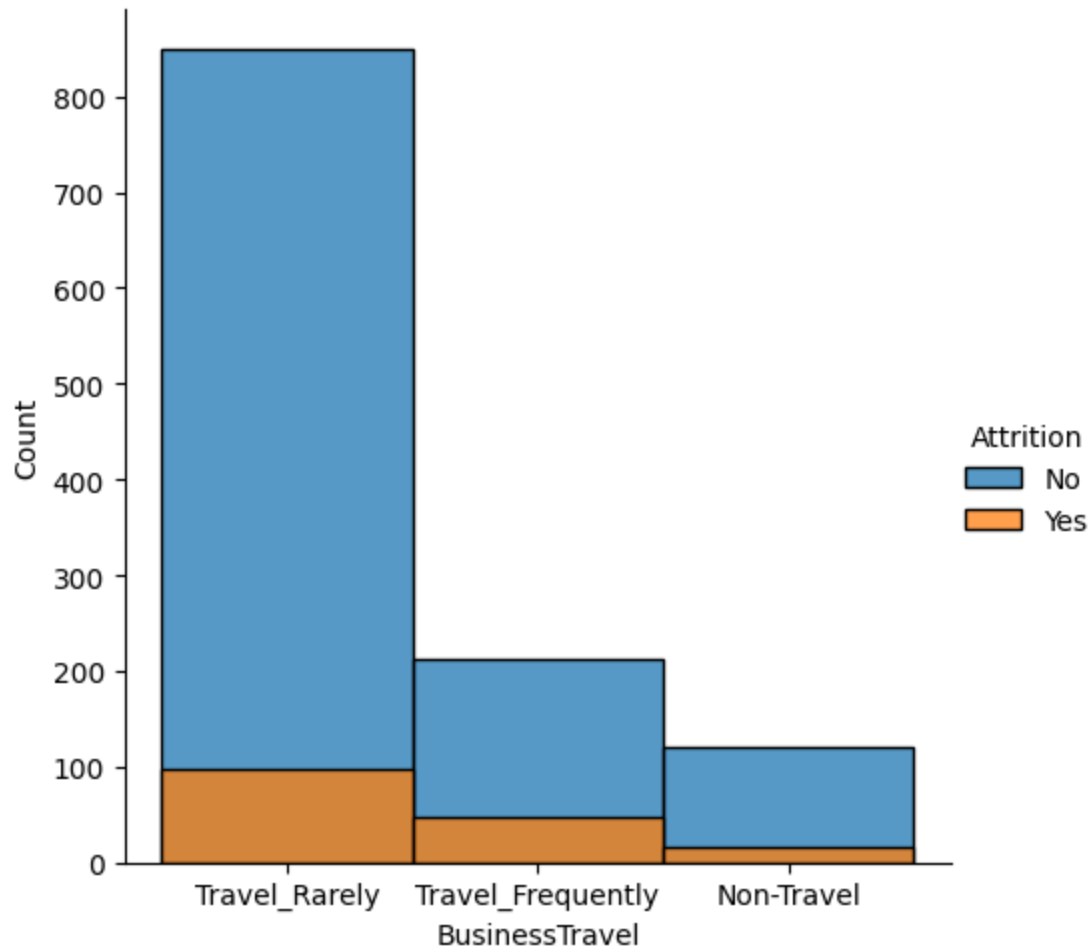
In [ ]: ### Initial dataset visualization with histograms
for c in data.columns:
    sns.FacetGrid(data,
                  hue="Attrition",
                  height= 5).map(sns.histplot,c).add_legend()


```

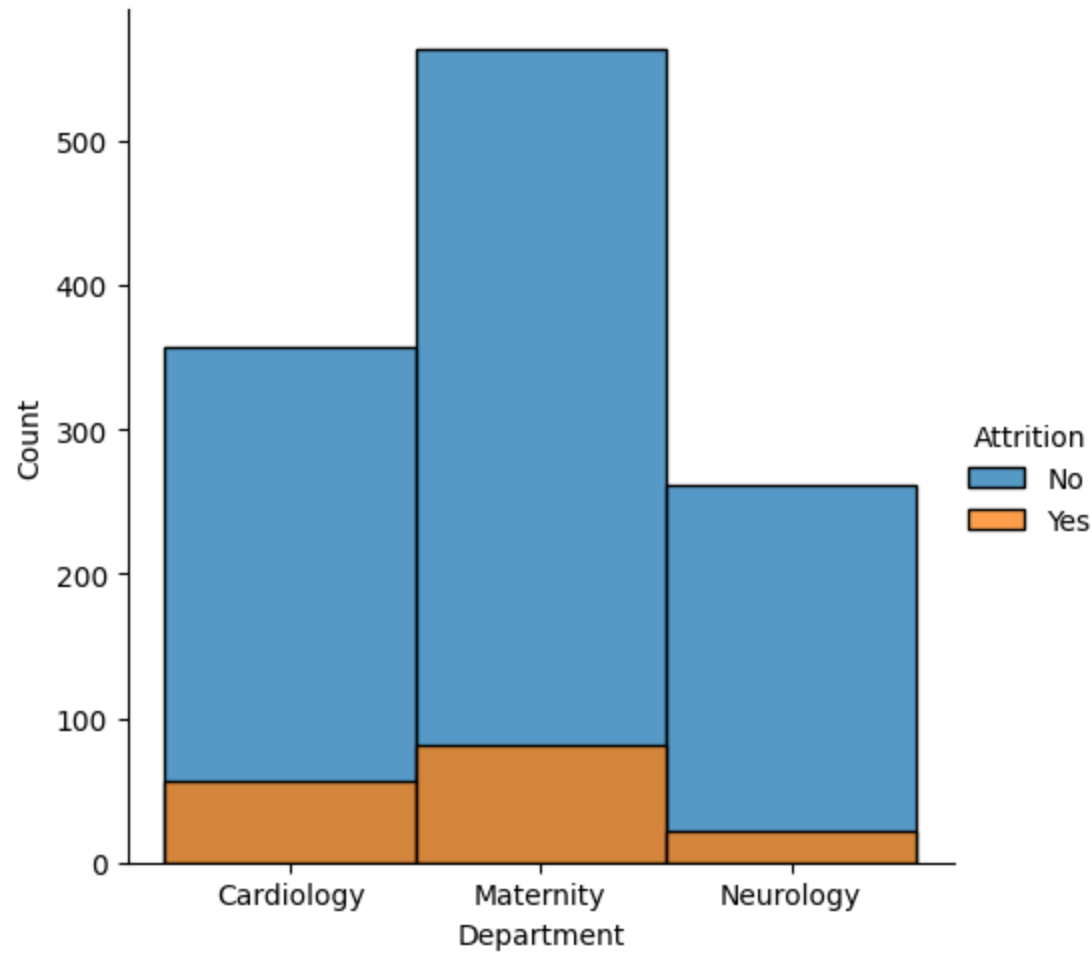








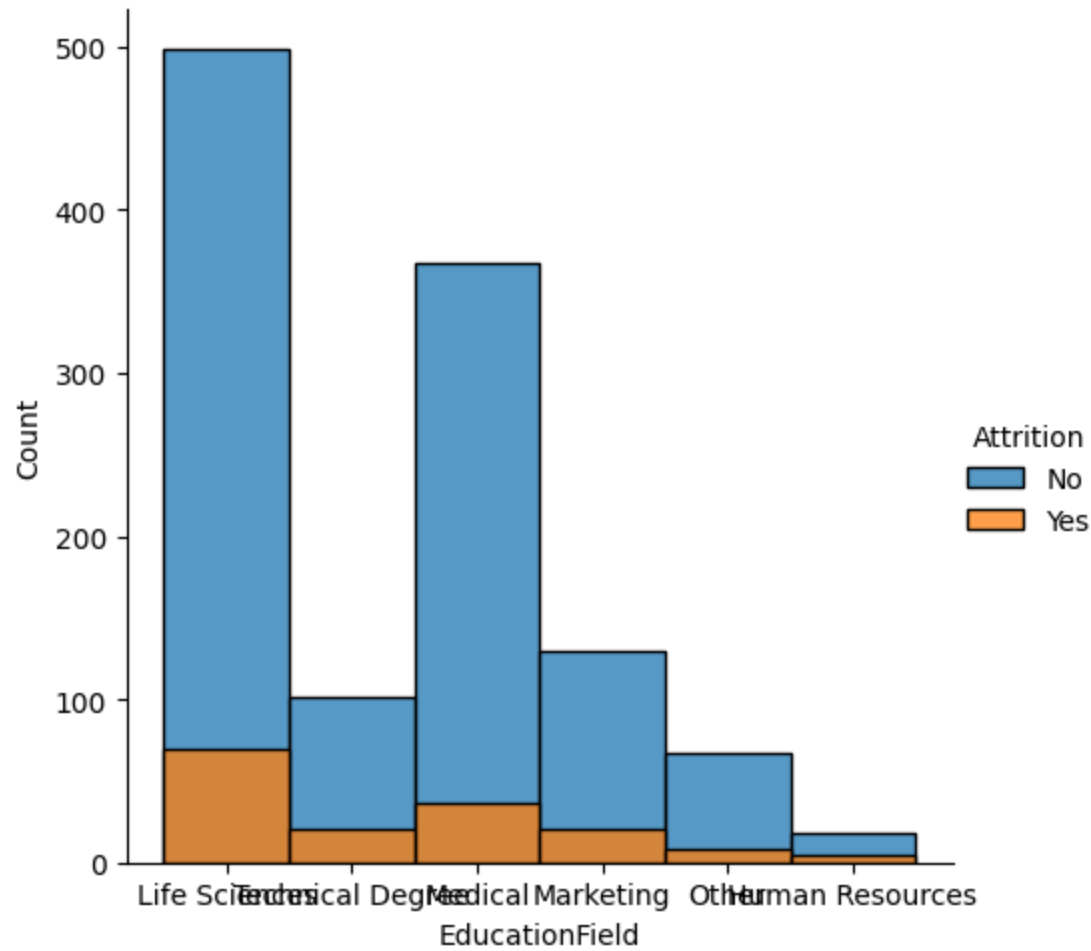




















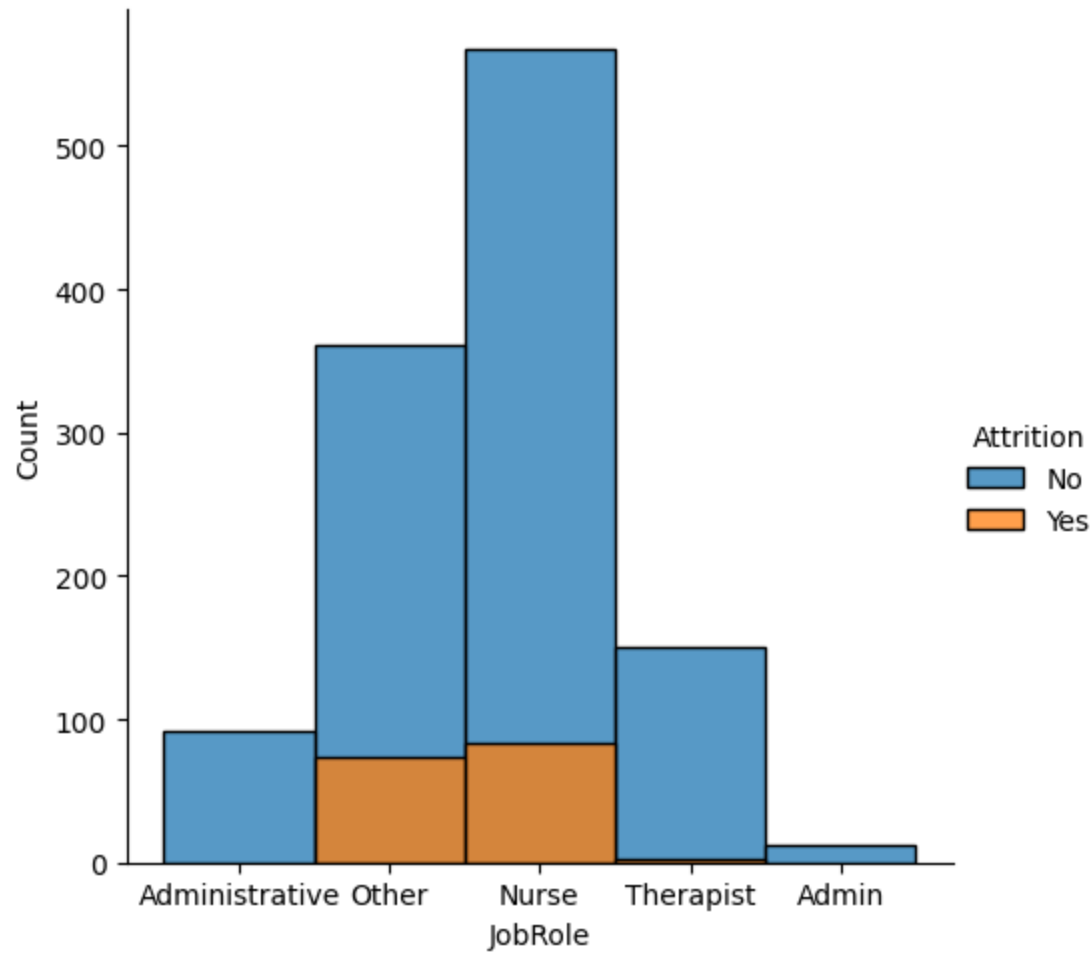





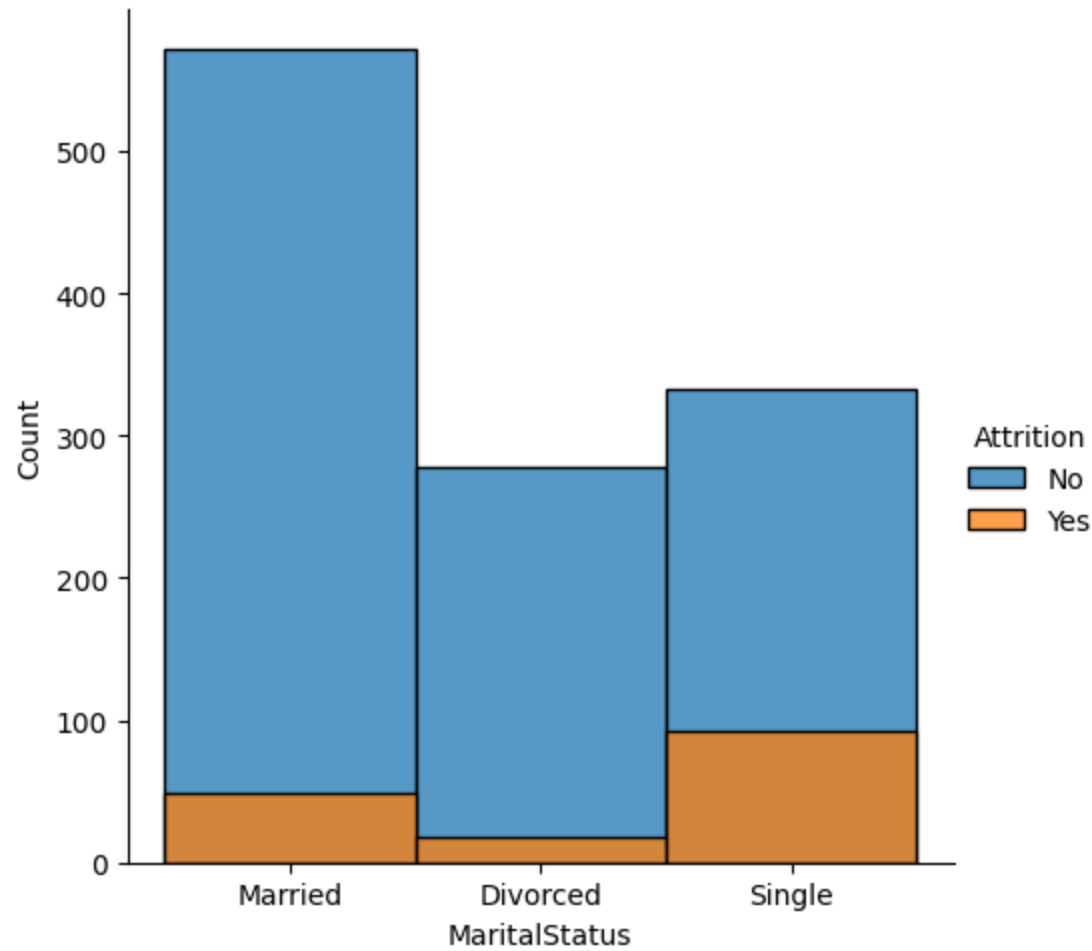





























































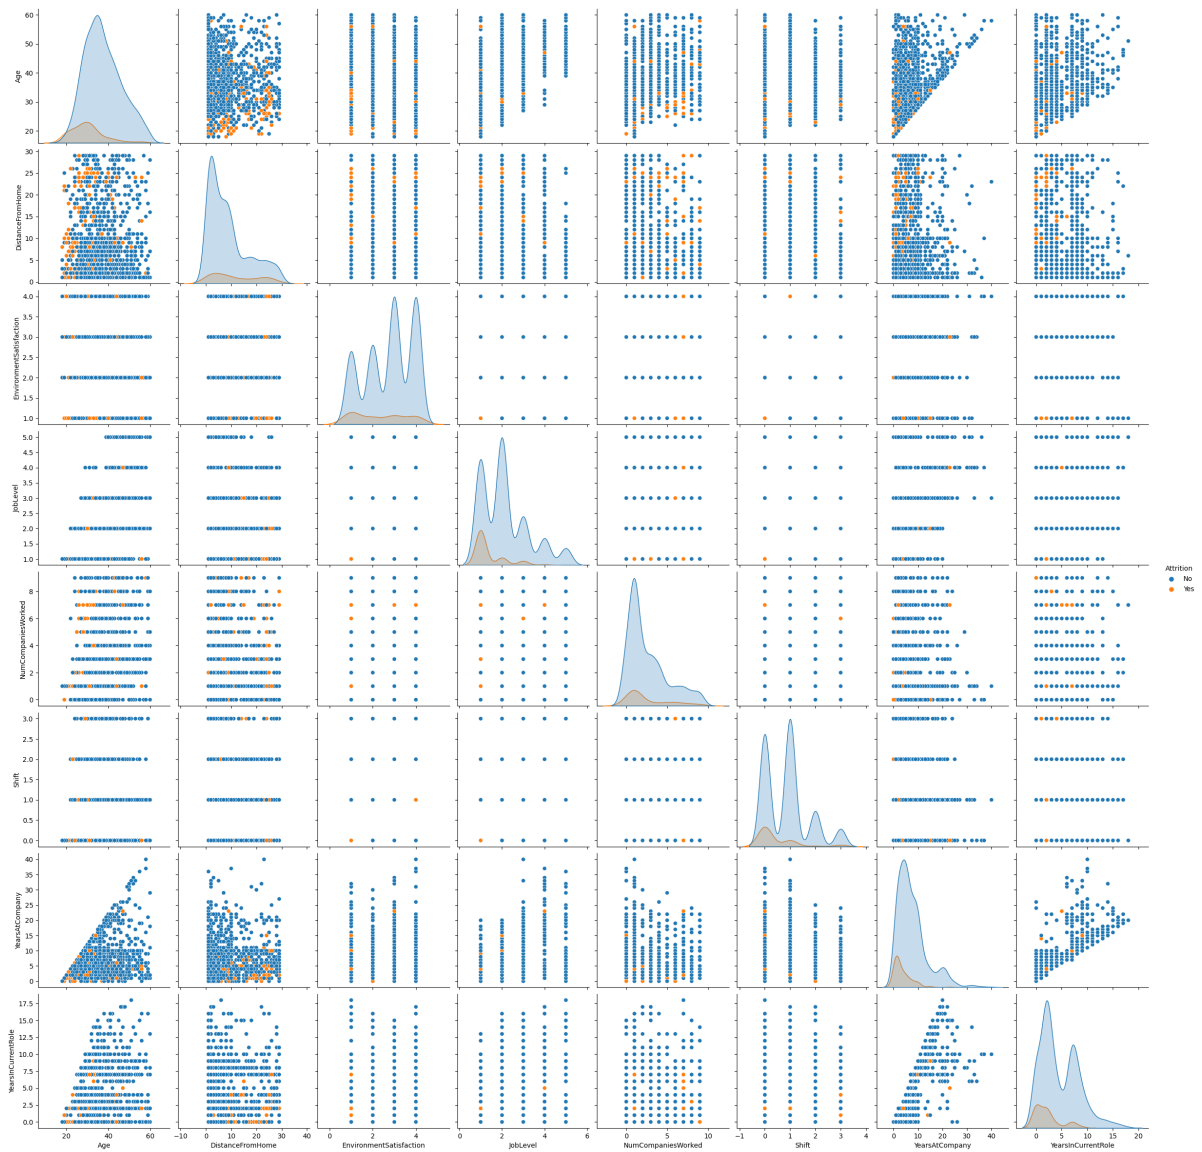


```
In [ ]: # Make pairwise scatter plots promising features seen in above analysis
promising_features = ["Age", "DistanceFromHome",
                      "EnvironmentSatisfaction",
                      "JobLevel", "JobRole",
                      "MaritalStatus",
                      "NumCompaniesWorked", "OverTime",
                      "Shift", "YearsAtCompany",
                      "YearsInCurrentRole",
                      "Attrition"]

promising_data = data[promising_features]

sns.pairplot(promising_data,
             hue="Attrition", height=3)
```

```
Out [ ]: <seaborn.axisgrid.PairGrid at 0x2abee75b0>
```



```
In [ ]: # Import other data
labels = data.pop("Attrition")
submission_data = pd.read_csv("test.csv")
print(data.shape)
print(submission_data.shape)
```

```
(1340, 34)
```

```
(336, 34)
```

```
In [ ]: # drop columns with only 1 unique value
for c in data.columns:
    print(c, " num unique values: ", len(data[c].unique()), "\n")
    if len(data[c].unique()) == 1:
        data.drop([c], axis=1, inplace=True)
        submission_data.drop([c], axis=1, inplace=True)
print(data.shape)
print(submission_data.shape)
```

EmployeeID num unique values: 1340

Age num unique values: 43

BusinessTravel num unique values: 3

DailyRate num unique values: 793

Department num unique values: 3

DistanceFromHome num unique values: 29

Education num unique values: 5

EducationField num unique values: 6

EmployeeCount num unique values: 1

EnvironmentSatisfaction num unique values: 4

Gender num unique values: 2

HourlyRate num unique values: 71

JobInvolvement num unique values: 4

JobLevel num unique values: 5

JobRole num unique values: 5

JobSatisfaction num unique values: 4

MaritalStatus num unique values: 3

MonthlyIncome num unique values: 1134

MonthlyRate num unique values: 1191

NumCompaniesWorked num unique values: 10

Over18 num unique values: 1

OverTime num unique values: 2

PercentSalaryHike num unique values: 15

PerformanceRating num unique values: 2

RelationshipSatisfaction num unique values: 4

StandardHours num unique values: 1

Shift num unique values: 4

TotalWorkingYears num unique values: 40

TrainingTimesLastYear num unique values: 7

WorkLifeBalance num unique values: 4

YearsAtCompany num unique values: 37

YearsInCurrentRole num unique values: 19

YearsSinceLastPromotion num unique values: 16

YearsWithCurrManager num unique values: 18

(1340, 31)

(336, 31)

See that there are 9 categorical columns which need to be converted to numerical.

See that there are many numerical columns which need to be binned by quantile.

Dropped columns with only 1 unique value -> can't get any information from those.

Can see that there are no null values, so we do not need to clean out rows or columns containing nulls.

```
In [ ]: # Binning
int_cols = []
for c in data.columns:
    if data[c].dtype == "int64":
        print(c, len(data[c].unique()))
        print(data[c].describe(), "\n\n")
```

EmployeeID 1340
count 1.340000e+03
mean 1.460265e+06
std 2.494821e+05
min 1.025177e+06
25% 1.237599e+06
50% 1.469862e+06
75% 1.670131e+06
max 1.886378e+06
Name: EmployeeID, dtype: float64

Age 43
count 1340.000000
mean 36.580597
std 9.013072
min 18.000000
25% 30.000000
50% 35.000000
75% 42.000000
max 60.000000
Name: Age, dtype: float64

DailyRate 793
count 1340.000000
mean 799.197761
std 399.333256
min 102.000000
25% 465.000000
50% 796.000000
75% 1153.000000
max 1499.000000
Name: DailyRate, dtype: float64

DistanceFromHome 29
count 1340.000000
mean 9.193284
std 8.141621
min 1.000000
25% 2.000000
50% 7.000000
75% 14.000000
max 29.000000
Name: DistanceFromHome, dtype: float64

Education 5
count 1340.000000
mean 2.924627
std 1.036088

```
min      1.000000
25%      2.000000
50%      3.000000
75%      4.000000
max      5.000000
Name: Education, dtype: float64
```

```
EnvironmentSatisfaction 4
count    1340.000000
mean      2.709701
std       1.099961
min       1.000000
25%       2.000000
50%       3.000000
75%       4.000000
max       4.000000
Name: EnvironmentSatisfaction, dtype: float64
```

```
HourlyRate 71
count    1340.000000
mean     65.559701
std      20.335025
min      30.000000
25%      48.000000
50%      65.000000
75%      83.000000
max     100.000000
Name: HourlyRate, dtype: float64
```

```
JobInvolvement 4
count    1340.000000
mean      2.717910
std       0.717523
min       1.000000
25%       2.000000
50%       3.000000
75%       3.000000
max       4.000000
Name: JobInvolvement, dtype: float64
```

```
JobLevel 5
count    1340.000000
mean      2.051493
std       1.104491
min       1.000000
25%       1.000000
50%       2.000000
75%       3.000000
```

```
max          5.000000
Name: JobLevel, dtype: float64
```

```
JobSatisfaction 4
count    1340.000000
mean      2.746269
std       1.111328
min       1.000000
25%       2.000000
50%       3.000000
75%       4.000000
max       4.000000
Name: JobSatisfaction, dtype: float64
```

```
MonthlyIncome 1134
count    1340.000000
mean     6433.381343
std      4687.058380
min      1051.000000
25%      2870.000000
50%      4876.500000
75%      8038.750000
max     19973.000000
Name: MonthlyIncome, dtype: float64
```

```
MonthlyRate 1191
count    1340.000000
mean     14290.377612
std      7166.995911
min      2094.000000
25%      7967.250000
50%     14288.500000
75%     20472.500000
max     26997.000000
Name: MonthlyRate, dtype: float64
```

```
NumCompaniesWorked 10
count    1340.000000
mean      2.600000
std       2.472794
min       0.000000
25%       1.000000
50%       1.000000
75%       4.000000
max       9.000000
Name: NumCompaniesWorked, dtype: float64
```

```
PercentSalaryHike 15
count      1340.000000
mean        15.168657
std         3.661956
min         11.000000
25%         12.000000
50%         14.000000
75%         18.000000
max         25.000000
Name: PercentSalaryHike, dtype: float64
```

```
PerformanceRating 2
count      1340.000000
mean         3.152239
std         0.359386
min         3.000000
25%         3.000000
50%         3.000000
75%         3.000000
max         4.000000
Name: PerformanceRating, dtype: float64
```

```
RelationshipSatisfaction 4
count      1340.000000
mean         2.700000
std         1.079858
min         1.000000
25%         2.000000
50%         3.000000
75%         4.000000
max         4.000000
Name: RelationshipSatisfaction, dtype: float64
```

```
Shift 4
count      1340.000000
mean         0.808209
std         0.856251
min         0.000000
25%         0.000000
50%         1.000000
75%         1.000000
max         3.000000
Name: Shift, dtype: float64
```

```
TotalWorkingYears 40
count      1340.000000
mean       11.222388
```

```
std          7.696043
min          0.000000
25%          6.000000
50%         10.000000
75%         15.000000
max          40.000000
Name: TotalWorkingYears, dtype: float64
```

```
TrainingTimesLastYear 7
count      1340.000000
mean        2.785821
std         1.263473
min         0.000000
25%         2.000000
50%         3.000000
75%         3.000000
max         6.000000
Name: TrainingTimesLastYear, dtype: float64
```

```
WorkLifeBalance 4
count      1340.000000
mean        2.771642
std         0.700007
min         1.000000
25%         2.000000
50%         3.000000
75%         3.000000
max         4.000000
Name: WorkLifeBalance, dtype: float64
```

```
YearsAtCompany 37
count      1340.000000
mean        7.070149
std         6.039663
min         0.000000
25%         3.000000
50%         5.000000
75%        10.000000
max        40.000000
Name: YearsAtCompany, dtype: float64
```

```
YearsInCurrentRole 19
count      1340.000000
mean        4.272388
std         3.677798
min         0.000000
25%         2.000000
50%         3.000000
```

```

75%          7.000000
max          18.000000
Name: YearsInCurrentRole, dtype: float64

```

```

YearsSinceLastPromotion 16
count      1340.000000
mean        2.175373
std         3.222376
min         0.000000
25%         0.000000
50%         1.000000
75%         3.000000
max        15.000000
Name: YearsSinceLastPromotion, dtype: float64

```

```

YearsWithCurrManager 18
count      1340.000000
mean        4.167164
std         3.581605
min         0.000000
25%         2.000000
50%         3.000000
75%         7.000000
max        17.000000
Name: YearsWithCurrManager, dtype: float64

```

```

In [ ]: # Binning
# found manually
cols_to_bin = ["Age", "DailyRate", "DistanceFromHome",
               "HourlyRate", "MonthlyIncome", "MonthlyRate",
               "PercentSalaryHike", "TotalWorkingYears",
               "YearsAtCompany", "YearsInCurrentRole",
               "YearsWithCurrManager", "NumCompaniesWorked",
               "YearsSinceLastPromotion",]

print(data.shape)
print(submission_data.shape)

# uneven 4 groups
for c in cols_to_bin:
    try:
        data[c+"_Even"] = pd.cut(data[c], 4, labels=False)
        submission_data[c+"_Even"] = pd.cut(submission_data[c], 4, labels=False)
    except:
        print("failed")
        pass

```

```
print(data.shape)
print(submission_data.shape)
```

```
(1340, 31)
(336, 31)
(1340, 44)
(336, 44)
```

```
In [ ]: # get dummies from int categorical data
```

```
# found manually
```

```
already_categorical = ["Education", "EnvironmentSatisfaction",
                       "JobInvolvement", "JobLevel", "JobSatisfaction",
                       "PerformanceRating", "RelationshipSatisfaction",
                       "Shift", "TrainingTimesLastYear",
                       "WorkLifeBalance"]
```

```
for c in already_categorical:
```

```
    temp_dummy = pd.get_dummies(data[c], prefix=c)
```

```
    data = pd.concat([data, temp_dummy], axis=1)
```

```
    sub_temp_dummy = pd.get_dummies(submission_data[c], prefix=c)
```

```
    submission_data = pd.concat([submission_data, sub_temp_dummy], axis=1)
```

```
print(data.shape)
```

```
print(submission_data.shape)
```

```
(1340, 87)
```

```
(336, 87)
```

```
In [ ]: # get dummies from obj categorical data
```

```
for c in data.columns:
```

```
    if data[c].dtype == "object":
```

```
        if len(data[c].unique()) == 2:
```

```
            data[c] = pd.factorize(data[c])[0]
```

```
            submission_data[c] = pd.factorize(submission_data[c])[0]
```

```
        else:
```

```
            temp_dummy = pd.get_dummies(data[c], prefix=c)
```

```
            data = pd.concat([data, temp_dummy], axis=1)
```

```
            data[c] = pd.factorize(data[c])[0]
```

```
            sub_temp_dummy = pd.get_dummies(submission_data[c], prefix=c)
```

```
            submission_data = pd.concat([submission_data, sub_temp_dummy], axis=1)
```

```
            submission_data[c] = pd.factorize(submission_data[c])[0]
```

```
print(data.shape)
```

```
print(submission_data.shape)
```

```
# turn label column to binary
```

```
labels = pd.DataFrame(pd.factorize(labels)[0], columns=["Attrition"])
```

```
(1340, 107)
```

```
(336, 107)
```


```
In [ ]: # Histograms for binned data and dummy data created
```


```
promising_data2 = data[data.columns[31:]]
```





```
promising_data2 = pd.concat([promising_data2, labels], axis=1)


for c in promising_data2.columns:
    sns.FacetGrid(promising_data2,
                  hue="Attrition",
                  height= 5).map(sns.histplot,c).add_legend()
```



































































































































































































































```
In [ ]: # Feature selection using coefficient weights
estimator = LogisticRegression(max_iter=120)
selector = RFECV(estimator, min_features_to_select=20,
                 step=1, cv=5)
selector = selector.fit(data, np.ravel(labels))
names = selector.get_feature_names_out()

# See which features are correlated with Attrition
print(data[names].columns)
print(data[names].shape)
```

```
Index(['BusinessTravel', 'Department', 'DistanceFromHome', 'Education',
      'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',
      'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus',
      'NumCompaniesWorked', 'OverTime', 'PercentSalaryHike',
      'PerformanceRating', 'RelationshipSatisfaction', 'Shift',
      'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',
      'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
      'YearsWithCurrManager', 'Age_Even', 'DailyRate_Even',
      'DistanceFromHome_Even', 'HourlyRate_Even', 'MonthlyIncome_Even',
      'MonthlyRate_Even', 'TotalWorkingYears_Even', 'YearsInCurrentRole_Ev
en',
      'YearsWithCurrManager_Even', 'NumCompaniesWorked_Even',
      'EnvironmentSatisfaction_1', 'EnvironmentSatisfaction_3',
      'EnvironmentSatisfaction_4', 'JobInvolvement_1', 'JobInvolvement_2',
      'JobInvolvement_3', 'JobInvolvement_4', 'JobLevel_1', 'JobLevel_2',
      'JobLevel_3', 'JobSatisfaction_1', 'JobSatisfaction_4',
      'RelationshipSatisfaction_4', 'Shift_0', 'Shift_1', 'Shift_2',
      'TrainingTimesLastYear_2', 'TrainingTimesLastYear_3',
      'WorkLifeBalance_1', 'WorkLifeBalance_3',
      'BusinessTravel_Travel_Frequently', 'BusinessTravel_Travel_Rarely',
      'Department_Cardiology', 'Department_Maternity',
      'EducationField_Marketing', 'EducationField_Medical',
      'JobRole_Therapist', 'MaritalStatus_Divorced', 'MaritalStatus_Marrie
d',
      'MaritalStatus_Single'],
      dtype='object')
(1340, 65)
```

```
In [ ]: # Eliminate redundant features
good_data = data[names]
selected_features = mrmr.mrmr_classif(good_data,
                                     np.ravel(labels),
                                     K=20)

print(selected_features)

uncorr_data = good_data[selected_features]
uncorr_sub_data = submission_data[selected_features]
```

```
100%|██████████| 20/20 [00:00<00:00, 86.50it/s]
['OverTime', 'JobLevel_1', 'BusinessTravel_Travel_Rarely', 'JobInvolvement',
 'Shift_0', 'WorkLifeBalance_1', 'Age_Even', 'DistanceFromHome_Even', 'EnvironmentSatisfaction',
 'YearsInCurrentRole', 'MaritalStatus_Single', 'TotalWorkingYears', 'JobSatisfaction',
 'JobLevel_2', 'JobInvolvement_1', 'YearsWithCurrManager', 'JobLevel', 'EnvironmentSatisfaction_1',
 'TrainingTimesLastYear_2', 'MaritalStatus']
```

```
In [ ]: # Save the selected data
fin_data = pd.concat([uncorr_data, labels], axis=1)
fin_data.to_csv("uncorr20_data.csv")
uncorr_sub_data.to_csv("uncorr20_sub_data.csv")
```

```
In [ ]: # Feature engineering by making polynomial features
poly = PolynomialFeatures(degree=2)

poly_data = poly.fit_transform(uncorr_data)
```

```
poly_sub = poly.fit_transform(uncorr_sub_data)

poly_names = poly.get_feature_names_out()
# print(poly_names)

# print(poly_data.shape)

poly_data = pd.DataFrame(poly_data, columns=poly_names)
poly_sub = pd.DataFrame(poly_sub, columns=poly_names)
# print(poly_data.head())
```

```
In [ ]: # Feature selection using coefficient weights
estimator = LogisticRegression(max_iter=120)
selector = RFECV(estimator, min_features_to_select=20,
                 step=1, cv=5)
selector = selector.fit(poly_data, np.ravel(labels))
names = selector.get_feature_names_out()

print(poly_data[names].columns)
print(poly_data[names].shape)
```

```

Index(['OverTime', 'YearsInCurrentRole', 'MaritalStatus_Single',
      'YearsWithCurrManager', 'OverTime^2', 'OverTime JobLevel_1',
      'OverTime YearsInCurrentRole', 'OverTime JobSatisfaction',
      'OverTime EnvironmentSatisfaction_1',
      'OverTime TrainingTimesLastYear_2', 'OverTime MaritalStatus',
      'JobLevel_1 WorkLifeBalance_1', 'JobLevel_1 Age_Even',
      'JobLevel_1 DistanceFromHome_Even', 'JobLevel_1 MaritalStatus_Singl
e',
      'JobLevel_1 EnvironmentSatisfaction_1',
      'BusinessTravel_Travel_Rarely WorkLifeBalance_1',
      'BusinessTravel_Travel_Rarely Age_Even',
      'BusinessTravel_Travel_Rarely DistanceFromHome_Even',
      'BusinessTravel_Travel_Rarely JobInvolvement_1',
      'BusinessTravel_Travel_Rarely JobLevel',
      'BusinessTravel_Travel_Rarely TrainingTimesLastYear_2',
      'BusinessTravel_Travel_Rarely MaritalStatus',
      'JobInvolvement MaritalStatus_Single', 'JobInvolvement JobSatisfacti
on',
      'JobInvolvement JobLevel_2', 'JobInvolvement EnvironmentSatisfaction
_1',
      'JobInvolvement MaritalStatus', 'Shift_0 WorkLifeBalance_1',
      'Shift_0 DistanceFromHome_Even', 'Shift_0 EnvironmentSatisfaction',
      'Shift_0 YearsInCurrentRole', 'Shift_0 MaritalStatus_Single',
      'Shift_0 JobInvolvement_1', 'Shift_0 YearsWithCurrManager',
      'Shift_0 EnvironmentSatisfaction_1', 'Shift_0 TrainingTimesLastYear_
2',
      'Shift_0 MaritalStatus', 'WorkLifeBalance_1 DistanceFromHome_Even',
      'WorkLifeBalance_1 YearsInCurrentRole',
      'WorkLifeBalance_1 TotalWorkingYears',
      'WorkLifeBalance_1 JobSatisfaction', 'WorkLifeBalance_1 JobLevel_2',
      'WorkLifeBalance_1 YearsWithCurrManager', 'Age_Even^2',
      'Age_Even EnvironmentSatisfaction', 'Age_Even YearsInCurrentRole',
      'Age_Even MaritalStatus_Single', 'Age_Even JobSatisfaction',
      'Age_Even JobLevel_2', 'Age_Even MaritalStatus',
      'DistanceFromHome_Even EnvironmentSatisfaction',
      'DistanceFromHome_Even JobLevel_2',
      'DistanceFromHome_Even JobInvolvement_1',
      'DistanceFromHome_Even TrainingTimesLastYear_2',
      'DistanceFromHome_Even MaritalStatus',
      'EnvironmentSatisfaction YearsInCurrentRole',
      'EnvironmentSatisfaction MaritalStatus_Single',
      'EnvironmentSatisfaction JobInvolvement_1',
      'EnvironmentSatisfaction YearsWithCurrManager',
      'EnvironmentSatisfaction TrainingTimesLastYear_2',
      'EnvironmentSatisfaction MaritalStatus',
      'YearsInCurrentRole MaritalStatus_Single',
      'YearsInCurrentRole MaritalStatus', 'MaritalStatus_Single^2',
      'MaritalStatus_Single TotalWorkingYears',
      'MaritalStatus_Single JobLevel_2',
      'MaritalStatus_Single JobInvolvement_1',
      'MaritalStatus_Single JobLevel',
      'MaritalStatus_Single EnvironmentSatisfaction_1',
      'MaritalStatus_Single MaritalStatus', 'JobSatisfaction JobLevel_2',
      'JobSatisfaction TrainingTimesLastYear_2',
      'JobSatisfaction MaritalStatus', 'JobLevel_2 JobInvolvement_1',
      'JobLevel_2 MaritalStatus', 'JobInvolvement_1 YearsWithCurrManager',

```



```

        'JobInvolvement_1 JobLevel', 'JobInvolvement_1 TrainingTimesLastYear_2',
        'JobInvolvement_1 MaritalStatus', 'YearsWithCurrManager JobLevel',
        'YearsWithCurrManager EnvironmentSatisfaction_1',
        'YearsWithCurrManager MaritalStatus',
        'EnvironmentSatisfaction_1 MaritalStatus', 'TrainingTimesLastYear_2^2',
        'MaritalStatus^2'],
        dtype='object')
(1340, 86)

```

```

In [ ]: # Eliminate redundant features
good_poly_data = poly_data[names]
selected_features = mrmr.mrmr_classif(good_poly_data,
                                     np.ravel(labels),
                                     K=20)

print(selected_features)

uncorr_poly_data = good_poly_data[selected_features]
uncorr_poly_sub_data = poly_sub[selected_features]

```

```
100%|██████████| 20/20 [00:00<00:00, 50.45it/s]
```

```


['OverTime JobLevel_1', 'Shift_0 EnvironmentSatisfaction_1', 'DistanceFromHome_Even TrainingTimesLastYear_2', 'JobInvolvement JobSatisfaction', 'OverTime MaritalStatus', 'BusinessTravel_Travel_Rarely Age_Even', 'JobLevel_1 MaritalStatus_Single', 'OverTime', 'Shift_0 JobInvolvement_1', 'YearsWithCurrManager', 'OverTime EnvironmentSatisfaction_1', 'JobLevel_1 WorkLifeBalance_1', 'OverTime^2', 'Age_Even EnvironmentSatisfaction', 'Shift_0 DistanceFromHome_Even', 'JobInvolvement JobLevel_2', 'OverTime TrainingTimesLastYear_2', 'YearsInCurrentRole', 'MaritalStatus_Single EnvironmentSatisfaction_1', 'OverTime JobSatisfaction']


```


```


In [ ]: # Generate histograms to check usefulness of features selected
fin_poly_data = pd.concat([uncorr_poly_data, labels], axis=1)
for c in fin_poly_data.columns:
    sns.FacetGrid(fin_poly_data,
                  hue="Attrition",
                  height= 5).map(sns.histplot,c).add_legend()


```



























































```
In [ ]: # save the final poly data
fin_poly_data = pd.concat([uncorr_poly_data, labels], axis=1)
fin_poly_data.to_csv("uncorr20_poly_data.csv")
uncorr_poly_sub_data.to_csv("uncorr20_poly_sub_data.csv")
```